

```
In [3]: import pandas as pd
```

```
In [4]: data=pd.read_csv("/home/placement/Desktop/reddy/fiat500.csv")
```

```
In [5]: data.describe()
```

```
Out[5]:
```

| | ID | engine_power | age_in_days | km | previous_owners | lat | lon | price |
|-------|-------------|--------------|-------------|---------------|-----------------|-------------|-------------|--------------|
| count | 1538.000000 | 1538.000000 | 1538.000000 | 1538.000000 | 1538.000000 | 1538.000000 | 1538.000000 | 1538.000000 |
| mean | 769.500000 | 51.904421 | 1650.980494 | 53396.011704 | 1.123537 | 43.541361 | 11.563428 | 8576.003901 |
| std | 444.126671 | 3.988023 | 1289.522278 | 40046.830723 | 0.416423 | 2.133518 | 2.328190 | 1939.958641 |
| min | 1.000000 | 51.000000 | 366.000000 | 1232.000000 | 1.000000 | 36.855839 | 7.245400 | 2500.000000 |
| 25% | 385.250000 | 51.000000 | 670.000000 | 20006.250000 | 1.000000 | 41.802990 | 9.505090 | 7122.500000 |
| 50% | 769.500000 | 51.000000 | 1035.000000 | 39031.000000 | 1.000000 | 44.394096 | 11.869260 | 9000.000000 |
| 75% | 1153.750000 | 51.000000 | 2616.000000 | 79667.750000 | 1.000000 | 45.467960 | 12.769040 | 10000.000000 |
| max | 1538.000000 | 77.000000 | 4658.000000 | 235000.000000 | 4.000000 | 46.795612 | 18.365520 | 11100.000000 |

```
In [6]: data.tail(10)#last 10 rows
```

```
Out[6]:
```

| | ID | model | engine_power | age_in_days | km | previous_owners | lat | lon | price |
|-------------|------|--------|--------------|-------------|--------|-----------------|-----------|----------|-------|
| 1528 | 1529 | lounge | 51 | 2861 | 126000 | 1 | 43.841980 | 10.51531 | 5500 |
| 1529 | 1530 | lounge | 51 | 731 | 22551 | 1 | 38.122070 | 13.36112 | 9900 |
| 1530 | 1531 | lounge | 51 | 670 | 29000 | 1 | 45.764648 | 8.99450 | 10800 |
| 1531 | 1532 | sport | 73 | 4505 | 127000 | 1 | 45.528511 | 9.59323 | 4750 |
| 1532 | 1533 | pop | 51 | 1917 | 52008 | 1 | 45.548000 | 11.54947 | 9900 |
| 1533 | 1534 | sport | 51 | 3712 | 115280 | 1 | 45.069679 | 7.70492 | 5200 |
| 1534 | 1535 | lounge | 74 | 3835 | 112000 | 1 | 45.845692 | 8.66687 | 4600 |
| 1535 | 1536 | pop | 51 | 2223 | 60457 | 1 | 45.481541 | 9.41348 | 7500 |
| 1536 | 1537 | lounge | 51 | 2557 | 80750 | 1 | 45.000702 | 7.68227 | 5990 |
| 1537 | 1538 | pop | 51 | 1766 | 54276 | 1 | 40.323410 | 17.56827 | 7900 |

```
In [7]: data1=data.drop(['ID','lat','lon'],axis=1)
```

```
In [8]: data1
```

```
Out[8]:
```

| | model | engine_power | age_in_days | km | previous_owners | price |
|------|--------|--------------|-------------|--------|-----------------|-------|
| 0 | lounge | 51 | 882 | 25000 | 1 | 8900 |
| 1 | pop | 51 | 1186 | 32500 | 1 | 8800 |
| 2 | sport | 74 | 4658 | 142228 | 1 | 4200 |
| 3 | lounge | 51 | 2739 | 160000 | 1 | 6000 |
| 4 | pop | 73 | 3074 | 106880 | 1 | 5700 |
| ... | ... | ... | ... | ... | ... | ... |
| 1533 | sport | 51 | 3712 | 115280 | 1 | 5200 |
| 1534 | lounge | 74 | 3835 | 112000 | 1 | 4600 |
| 1535 | pop | 51 | 2223 | 60457 | 1 | 7500 |
| 1536 | lounge | 51 | 2557 | 80750 | 1 | 5990 |
| 1537 | pop | 51 | 1766 | 54276 | 1 | 7900 |

1538 rows × 6 columns

```
In [9]: data1['model']=data1['model'].map({'lounge':1,'pop':2,'sport':3})
```

```
In [10]: data1
```

```
Out[10]:
```

| | model | engine_power | age_in_days | km | previous_owners | price |
|------|-------|--------------|-------------|--------|-----------------|-------|
| 0 | 1 | 51 | 882 | 25000 | 1 | 8900 |
| 1 | 2 | 51 | 1186 | 32500 | 1 | 8800 |
| 2 | 3 | 74 | 4658 | 142228 | 1 | 4200 |
| 3 | 1 | 51 | 2739 | 160000 | 1 | 6000 |
| 4 | 2 | 73 | 3074 | 106880 | 1 | 5700 |
| ... | ... | ... | ... | ... | ... | ... |
| 1533 | 3 | 51 | 3712 | 115280 | 1 | 5200 |
| 1534 | 1 | 74 | 3835 | 112000 | 1 | 4600 |
| 1535 | 2 | 51 | 2223 | 60457 | 1 | 7500 |
| 1536 | 1 | 51 | 2557 | 80750 | 1 | 5990 |
| 1537 | 2 | 51 | 1766 | 54276 | 1 | 7900 |

1538 rows × 6 columns

```
In [11]: data2=data1.loc[(data.model=='lounge')]
```

```
In [12]: data2
```

```
Out[12]:
```

| | model | engine_power | age_in_days | km | previous_owners | price |
|------|-------|--------------|-------------|--------|-----------------|-------|
| 0 | 1 | 51 | 882 | 25000 | 1 | 8900 |
| 3 | 1 | 51 | 2739 | 160000 | 1 | 6000 |
| 6 | 1 | 51 | 731 | 11600 | 1 | 10750 |
| 7 | 1 | 51 | 1521 | 49076 | 1 | 9190 |
| 11 | 1 | 51 | 366 | 17500 | 1 | 10990 |
| ... | ... | ... | ... | ... | ... | ... |
| 1528 | 1 | 51 | 2861 | 126000 | 1 | 5500 |
| 1529 | 1 | 51 | 731 | 22551 | 1 | 9900 |
| 1530 | 1 | 51 | 670 | 29000 | 1 | 10800 |
| 1534 | 1 | 74 | 3835 | 112000 | 1 | 4600 |
| 1536 | 1 | 51 | 2557 | 80750 | 1 | 5990 |

1094 rows × 6 columns

```
In [13]: y=data2['price']  
x=data2.drop('price',axis=1)
```

In [14]:

y

Out[14]:

| | |
|------|-------|
| 0 | 8900 |
| 3 | 6000 |
| 6 | 10750 |
| 7 | 9190 |
| 11 | 10990 |
| | ... |
| 1528 | 5500 |
| 1529 | 9900 |
| 1530 | 10800 |
| 1534 | 4600 |
| 1536 | 5990 |

Name: price, Length: 1094, dtype: int64

In [15]:

x

Out[15]:

| | model | engine_power | age_in_days | km | previous_owners |
|------|-------|--------------|-------------|--------|-----------------|
| 0 | 1 | 51 | 882 | 25000 | 1 |
| 3 | 1 | 51 | 2739 | 160000 | 1 |
| 6 | 1 | 51 | 731 | 11600 | 1 |
| 7 | 1 | 51 | 1521 | 49076 | 1 |
| 11 | 1 | 51 | 366 | 17500 | 1 |
| ... | ... | ... | ... | ... | ... |
| 1528 | 1 | 51 | 2861 | 126000 | 1 |
| 1529 | 1 | 51 | 731 | 22551 | 1 |
| 1530 | 1 | 51 | 670 | 29000 | 1 |
| 1534 | 1 | 74 | 3835 | 112000 | 1 |
| 1536 | 1 | 51 | 2557 | 80750 | 1 |

1094 rows × 5 columns

```
In [16]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=42)
```

```
In [17]: x_test.head(5)
```

```
Out[17]:
```

| | model | engine_power | age_in_days | km | previous_owners |
|------|-------|--------------|-------------|--------|-----------------|
| 676 | 1 | 51 | 762 | 18609 | 1 |
| 215 | 1 | 51 | 701 | 25000 | 1 |
| 146 | 1 | 51 | 4018 | 152900 | 1 |
| 1319 | 1 | 51 | 731 | 20025 | 1 |
| 1041 | 1 | 51 | 640 | 38231 | 1 |

```
In [18]: x_train.shape
```

```
Out[18]: (732, 5)
```

```
In [19]: y_train.shape
```

```
Out[19]: (732,)
```

```
In [20]: x_train.head()
```

```
Out[20]:
```

| | model | engine_power | age_in_days | km | previous_owners |
|------|-------|--------------|-------------|-------|-----------------|
| 441 | 1 | 51 | 762 | 36448 | 1 |
| 701 | 1 | 51 | 701 | 27100 | 1 |
| 695 | 1 | 51 | 3197 | 51083 | 1 |
| 1415 | 1 | 51 | 670 | 33000 | 1 |
| 404 | 1 | 51 | 456 | 14000 | 1 |

```
In [21]: y_train.head()
```

```
Out[21]: 441      8980
         701     10300
         695      5880
        1415     10490
         404      9499
         Name: price, dtype: int64
```

```
In [22]: x_test.head()
```

```
Out[22]:
```

| | model | engine_power | age_in_days | km | previous_owners |
|------|-------|--------------|-------------|--------|-----------------|
| 676 | 1 | 51 | 762 | 18609 | 1 |
| 215 | 1 | 51 | 701 | 25000 | 1 |
| 146 | 1 | 51 | 4018 | 152900 | 1 |
| 1319 | 1 | 51 | 731 | 20025 | 1 |
| 1041 | 1 | 51 | 640 | 38231 | 1 |

```
In [23]: y_test.head()
```

```
Out[23]: 676      10250
         215      9790
         146      5500
        1319      9900
        1041      8900
         Name: price, dtype: int64
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```


In []:

In []:

```
In [33]: from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import Ridge
#ridge regression
alpha = [1e-15, 1e-10, 1e-8, 1e-4, 1e-3, 1e-2, 1, 5, 10, 20, 30]

ridge = Ridge()

parameters = {'alpha': alpha}

ridge_regressor = GridSearchCV(ridge, parameters)

ridge_regressor.fit(x_train, y_train)
```

```
Out[33]: GridSearchCV(estimator=Ridge(),
                      param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.0001, 0.001, 0.01, 1,
                                             5, 10, 20, 30]})
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [25]: ridge_regressor.best_params_
```

```
Out[25]: {'alpha': 30}
```

```
In [26]: ridge=Ridge(alpha=30)
ridge.fit(x_train,y_train)
y_pred_ridge=ridge.predict(x_test)
```

```
In [27]: from sklearn.metrics import mean_squared_error
Ridge_Error=mean_squared_error(y_pred_ridge,y_test)
Ridge_Error
```

```
Out[27]: 519771.8129989742
```

```
In [28]: from sklearn.metrics import r2_score  
r2_score(y_test,y_pred_ridge)
```

```
Out[28]: 0.8373030813683995
```


```
In [29]: Results=pd.DataFrame(columns=['Actual','predicted'])  
Results['Actual']=y_test  
Results['predicted']=y_pred_ridge  
Results=Results.reset_index()  
Results['ID']=Results.index  
Results.head(10)
```

```
Out[29]:
```

| | index | Actual | predicted | ID |
|---|-------|--------|--------------|----|
| 0 | 676 | 10250 | 10045.347779 | 0 |
| 1 | 215 | 9790 | 9989.171535 | 1 |
| 2 | 146 | 5500 | 4769.099603 | 2 |
| 3 | 1319 | 9900 | 10048.683238 | 3 |
| 4 | 1041 | 8900 | 9813.944798 | 4 |
| 5 | 1425 | 9500 | 8678.143561 | 5 |
| 6 | 409 | 10450 | 10173.797921 | 6 |
| 7 | 617 | 9790 | 10180.627008 | 7 |
| 8 | 1526 | 9300 | 9107.315259 | 8 |
| 9 | 1010 | 4600 | 5625.007407 | 9 |


```
In [30]: import seaborn as sns
import matplotlib.pyplot as plt
sns.lineplot(x='ID', y='Actual', data=Results.head(50))
sns.lineplot(x='ID', y='predicted', data=Results.head(50))
plt.plot()
```

Out[30]: []



```
In [31]: import seaborn as sns
import matplotlib.pyplot as plt
sns.lineplot(x='ID',y='Actual',data=Results.tail(50))
sns.lineplot(x='ID',y='predicted',data=Results.tail(50))
plt.plot()
```

Out[31]: []



In []:

